Fraud Detection Using Azure Data Engineering

Project Overview

This project is designed to detect potentially fraudulent transactions based on the distance between a customer's registered city and the city where the transaction occurred. Fraud criticality is categorized as **Low, Medium, or High** based on the calculated distance. The data pipeline is built using **Azure Data Factory (ADF), Azure Data Lake Storage (ADLS), Azure Databricks, and Azure SQL Database**.

Requirement

Based on the city in which the transaction has happened, the fraud criticality needs to be calculated based on the below conditions,

- **Low:** Distance < 500 km
- **Medium:** Distance 5000-1000 km
- **High:** Distance > 1000 km

The following reports need to be developed based on the Fraud criticality,

- Bar chart displaying Fraud Alert Counts for past 12 months.
- Pie chart showing Fraud criticality trend
- Top 10 Customers and Cities of High Fraud Criticality
- Pie chart showing Age Group criticality trend

Tech Stack

- **Azure Data Factory (ADF)** – Data ingestion and orchestration.
- **Azure Data Lake Storage (ADLS)** – Raw and landing data storage.
- **Azure Databricks** – Data transformation and fraud analysis.
- **Azure SQL Database** – Storing transformed data for reporting.
- **Power BI** – Visualization and reporting.

Source Datasets Used

1. **Customer Metadata** – Contains customer details, including city.

Customer	Name	Email	CityID	BirthYear	Gender	SignupDate
C001	Umang Chokshi	inaaya-36@tak.com	8	1998	Female	31-07-2021
C002	Kartik Dada	keyachaudhuri@hotmail.com	10	1976	Female	06-04-2024
C003	Tushar Kar	pari98@yahoo.com	9	1993	Female	24-07-2021

2. **City Distance Metadata** – Contains distance between every two cities in India.

CityID_1	CityName_1	CityID_2	CityName_2	Distance_km
1	Mumbai	1	Mumbai	0
1	Mumbai	2	Delhi	852
1	Mumbai	3	Bangalore	1940

3. **Transaction Data** – Contains transaction details, including transaction city.

TransactionID	Customer	CityID	Date	Amount	PaymentType	Status
T0001	C071	6	12-02-2025	8792.32	Debit Card	Success
T0002	C029	5	11-02-2025	40834.75	UPI	Pending
T0003	C007	7	08-02-2025	37359.97	UPI	Failed
T0004	C014	2	08-02-2025	17375.03	Net Banking	Success

Data Model

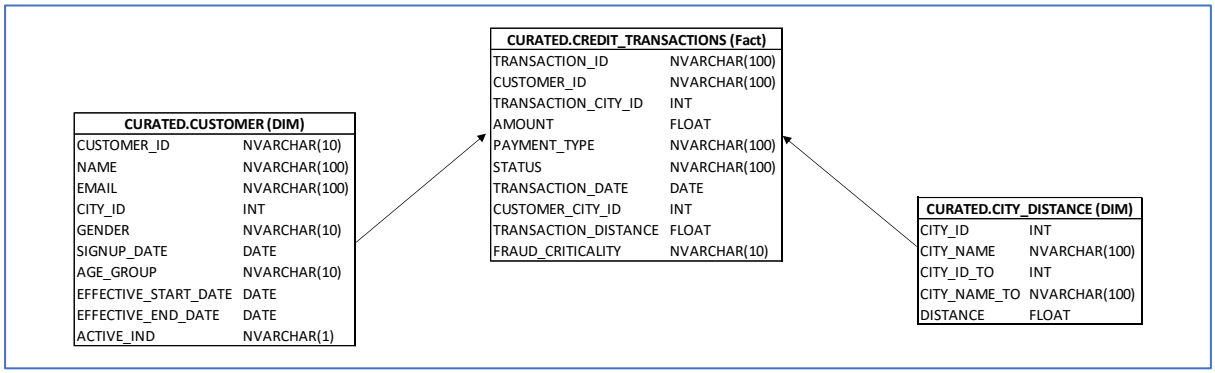

Architecture

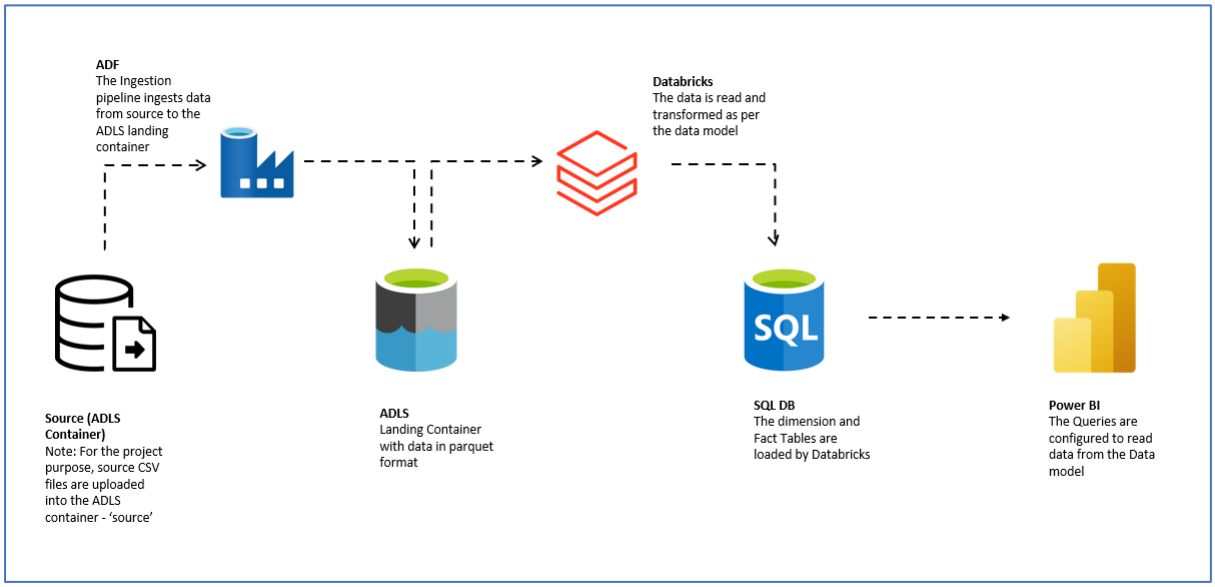

Detailed Implementation:

Azure Data Lake Storage:

Two containers are created,

Source – to store the raw csv source files.

Landing – to store the parquet file ingested by ADF

NOTE: This action is implemented to mimic the data ingestion from source to ADLS

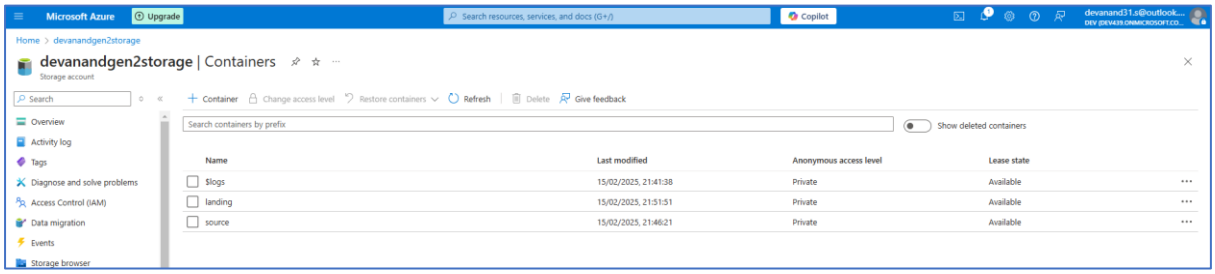

The transaction data files are stored in YYYY/MM/DD date folder format.

Azure Data Factory:

The ADF contains a main pipeline which starts with setting the load date (current date -1 day). It then executes two ingestion pipeline (Master data & Transaction data) and one curation pipeline.

PL_FRAUD_TRANSACTION_REPORT (main pipeline):

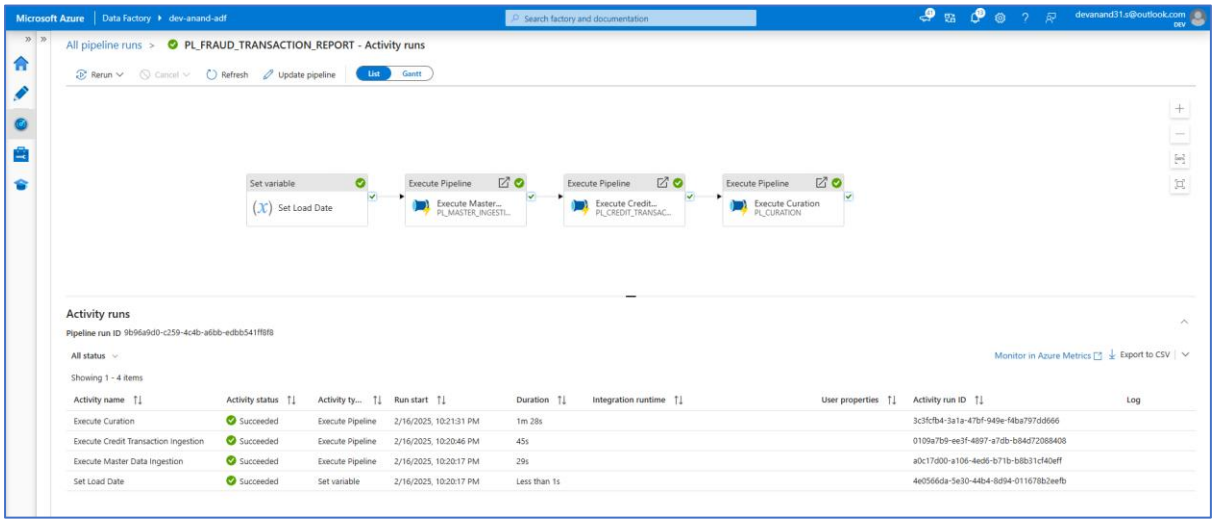

PL_MASTER_INGESTION: Ingests Customer and Cities data into landing container on full load basis

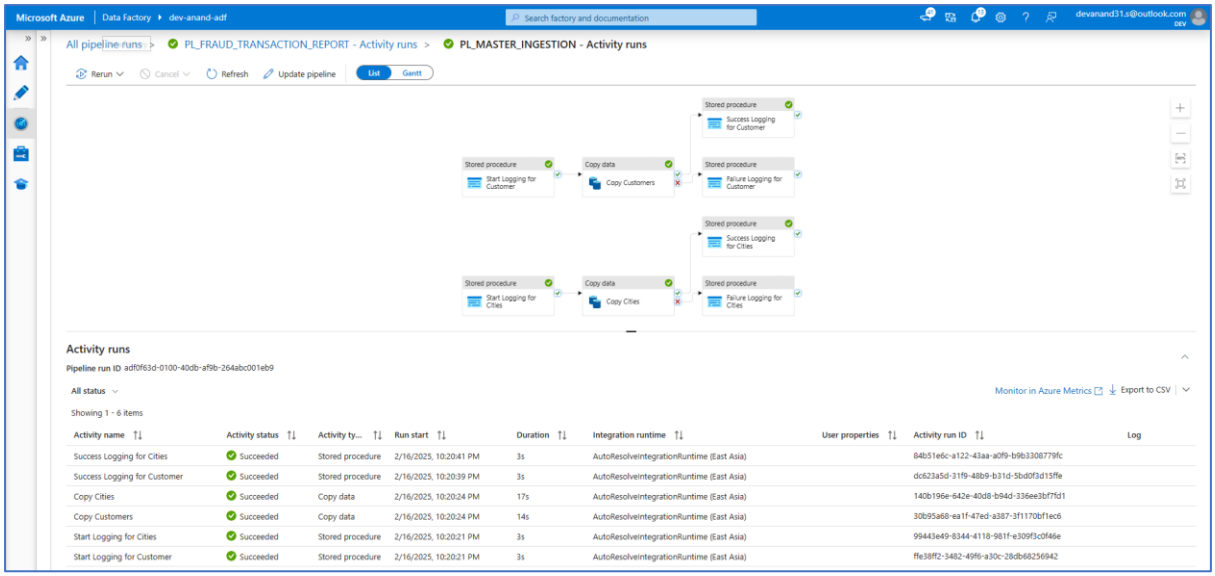

PL_CREDIT_TRANSACTION_INGESTION: For ingesting the transaction data on an incremental basis

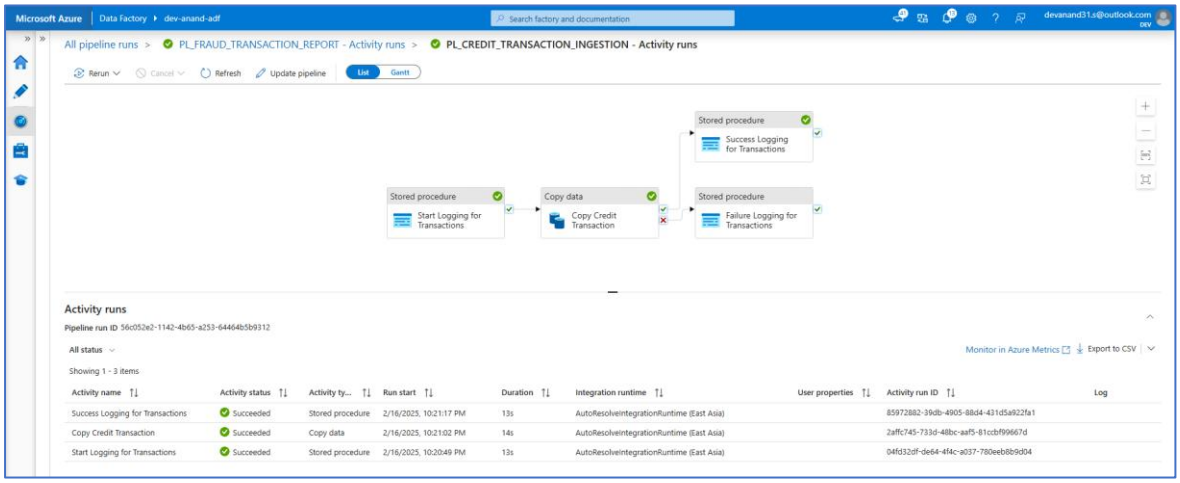

Incremental logic: The Load date is passed as parameter from the main pipeline which is used in the wild card path

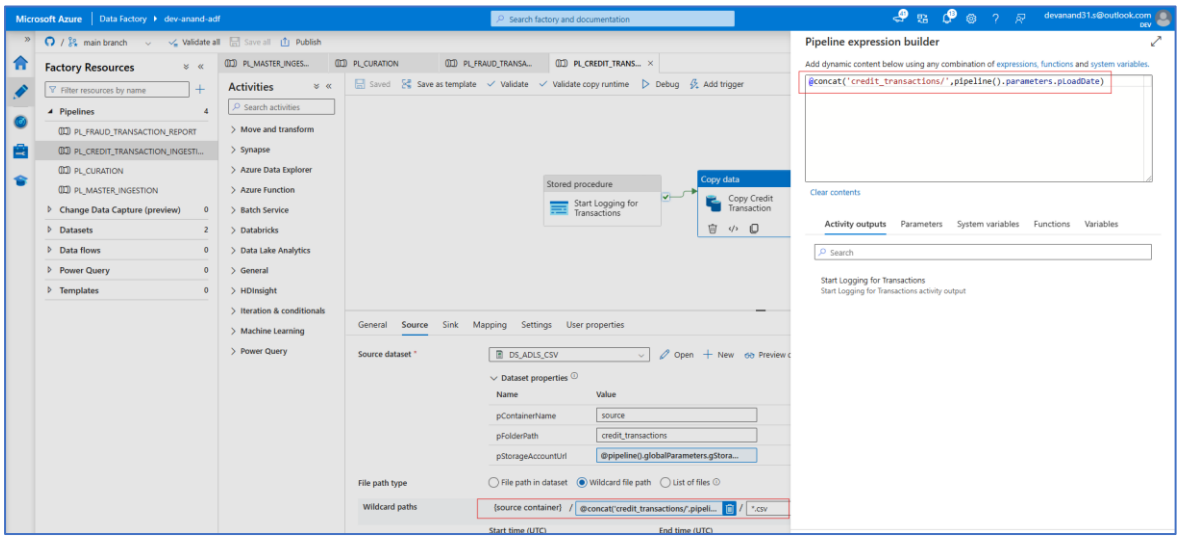

PL_CURATION: The curation pipeline contains two notebook activity. One for executing the master curation and another on for executing the transaction curation

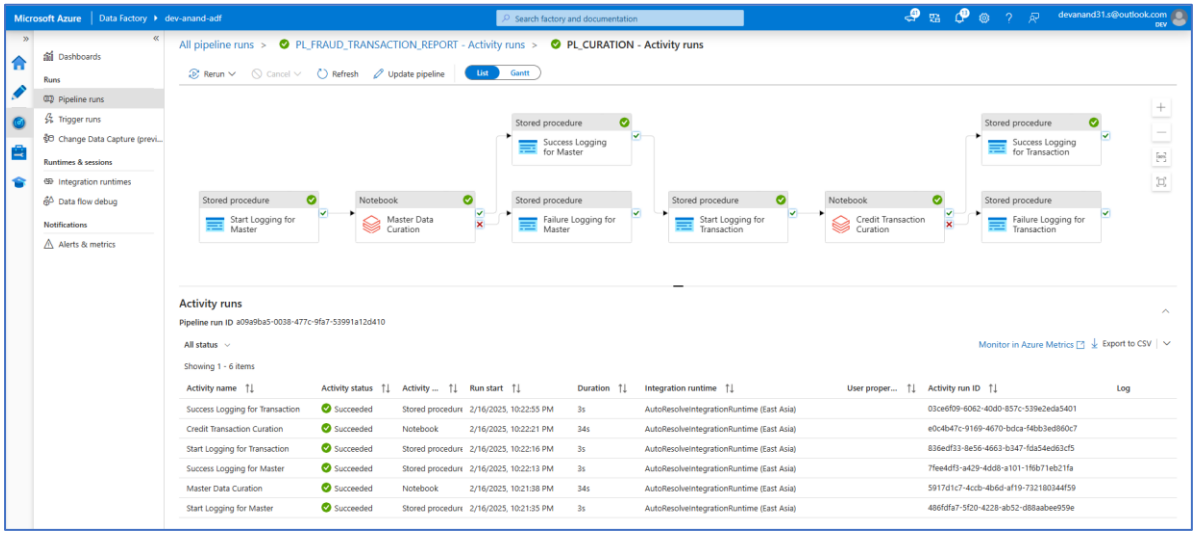

All pipelines are implemented with logging logic in which the stored procedure is called. This stored procedure INSERTs and UPDATEs the necessary logs at the appropriate steps.

Logging Table:

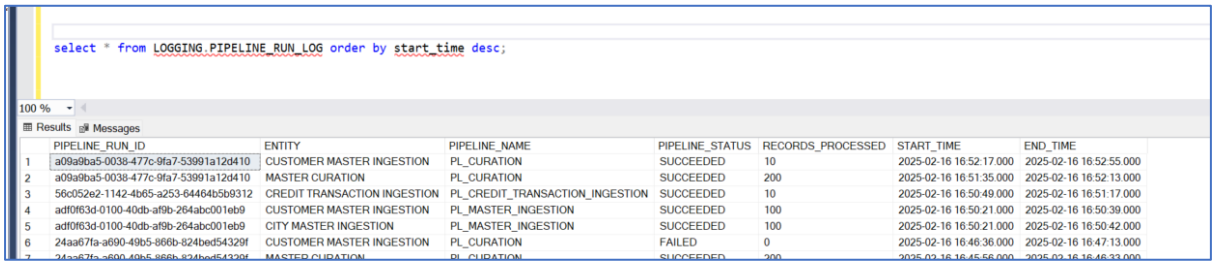

```
select * from LOGGING.PIPELINE_RUN_LOG order by start_time desc;
```

100 %

Results | Messages

	PIPELINE_RUN_ID	ENTITY	PIPELINE_NAME	PIPELINE_STATUS	RECORDS_PROCESSED	START_TIME	END_TIME
1	a09a9ba5-0038-477c-9fa7-53991a12d410	CUSTOMER MASTER INGESTION	PL_CURATION	SUCCEEDED	10	2025-02-16 16:52:17.000	2025-02-16 16:52:55.000
2	a09a9ba5-0038-477c-9fa7-53991a12d410	MASTER CURATION	PL_CURATION	SUCCEEDED	200	2025-02-16 16:51:35.000	2025-02-16 16:52:13.000
3	56c052e2-1142-4b65-a253-64464b5b9312	CREDIT TRANSACTION INGESTION	PL_CREDIT_TRANSACTION_INGESTION	SUCCEEDED	10	2025-02-16 16:50:49.000	2025-02-16 16:51:17.000
4	adf0f63d-0100-40db-af9b-264abc001eb9	CUSTOMER MASTER INGESTION	PL_MASTER_INGESTION	SUCCEEDED	100	2025-02-16 16:50:21.000	2025-02-16 16:50:39.000
5	adf0f63d-0100-40db-af9b-264abc001eb9	CITY MASTER INGESTION	PL_MASTER_INGESTION	SUCCEEDED	100	2025-02-16 16:50:21.000	2025-02-16 16:50:42.000
6	24aa67fa-a690-49b5-866b-824bed54329f	CUSTOMER MASTER INGESTION	PL_CURATION	FAILED	0	2025-02-16 16:46:36.000	2025-02-16 16:47:13.000

Databricks:

Connection and Helper Function:

- This notebook contains the necessary properties declared for establishing connections from Databricks notebook to ADLS and SQL DB.
- This notebook also contains the helped function that will accept the query string. This query string is executed in the SQL DB through ODBC Driver 17 for SQL Server.

Master Table Curation

- The customers and cities files from landing container are read, processed and loaded to the CURATED.CUSTOMER and CURATED.CITY_DISTANCE tables.
- The cities data are overwritten every time the notebook is executed.
- From the customer data, the age group of each customer is calculated. Every time when the notebook is executed, the data is compared against existing data in the table and instead of rewriting the updated data, the old data is marked as Inactive and the new updated data is inserted as Active. This helps in maintaining the customer history (SCD2) for future analysis.

Transaction Curation:

- This notebook reads the incremental transactions on a (current date – 1 day) basis.
- Using the customer and city curated data, the distance between the customer city and the city where the transaction has happened is calculated. Based on this distance, the criticality for fraud is calculated.
    - **Low:** Distance < 500 km
    - **Medium:** Distance 5000-1000 km
    - **High:** Distance > 1000 km

SQL Database:

- CURATED.CREDIT_TRANSACTIONS – This is the fact tables loaded with transaction containing the calculated fraud criticality column. It gets loaded on an incremental basis.
- CURATED.CUSTOMER – Contains customer master data. Follows SCD2.
- CURATED.CITY_DISTANCE – Contains city master data and is loaded on full load basis.
- LOGGING.SP_PIPELINE_RUN_LOG_START, LOGGING.SP_PIPELINE_RUN_LOG_END – Stored Procedures used to INSERT and UPDATE logs respectively
- LOGGING.PIPELINE_RUN_LOG – Contains pipeline execution logs

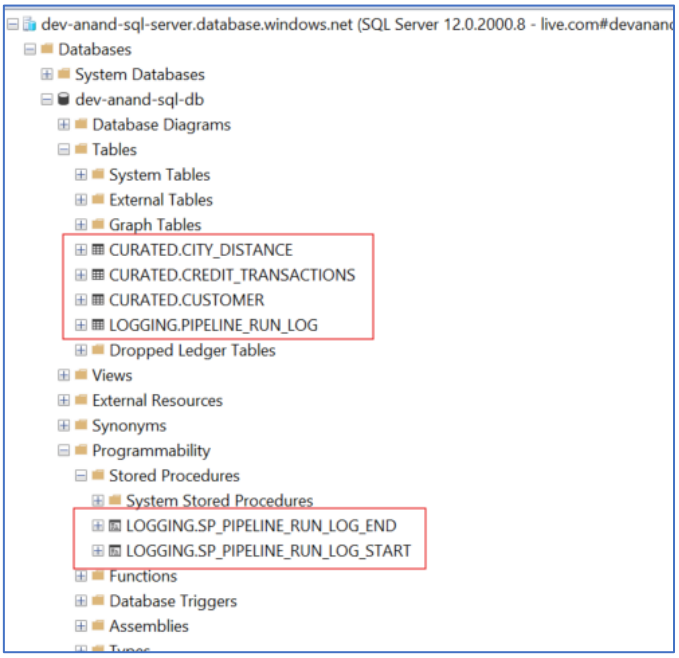

Power BI:

Finally, in the Power BI, 5 different queries are created to generate reports as per the requirement.

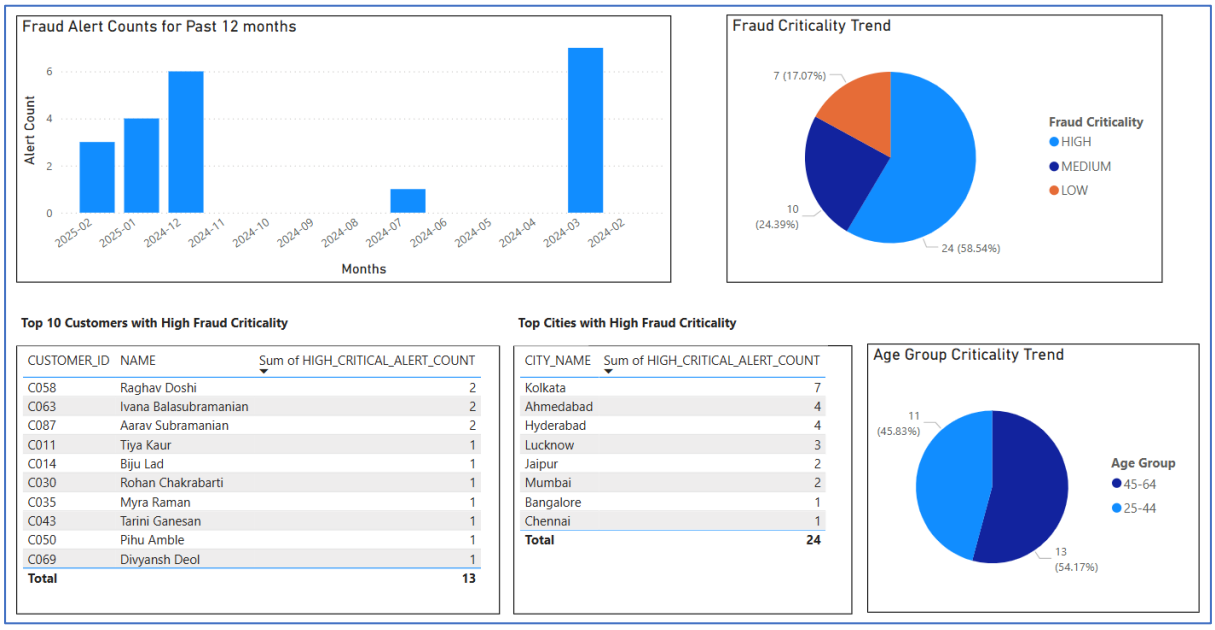

Security:

ADF <-> ADLS:

The Connection for the ADF to read and write files from and to ADLS is performed using managed identity with below access roles.

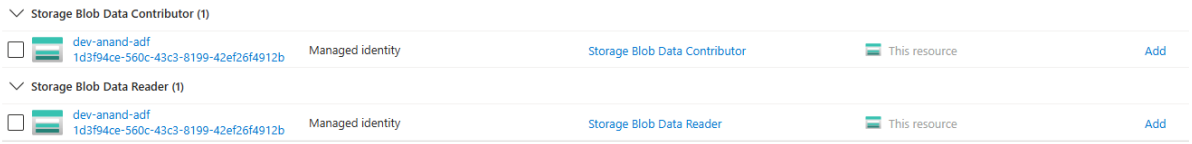

ADF <-> SQL DB:

The Connection for the ADF to read and write files from and to SQL DB is performed using managed identity with below access roles.

Databricks with ADF, ADLS and SQL DB:

Since this is a personal account, only the Standard edition of Databricks workspace is created. This version of Databricks does not allow Managed identity. So, all the Databricks related access are established used access tokens/access keys via Key Vault.

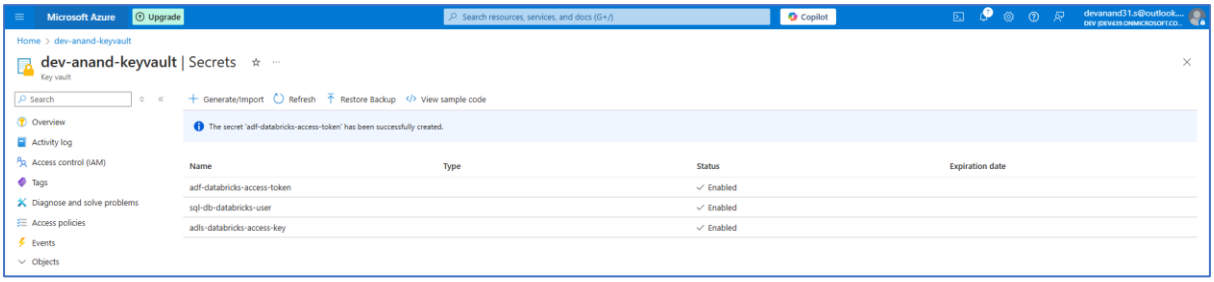

Power BI <-> SQL DB

A SQL user – powerbi_user is created for the Power BI to access SQL DB.

```
USE MASTER
GO
CREATE LOGIN powerbi_user WITH PASSWORD = '<<encrypted>>'
GO
CREATE USER powerbi_user FOR LOGIN powerbi_user
GO
ALTER ROLE db_datareader ADD MEMBER powerbi_user
GO
```